

Ανάπτυξη Λογισμικού για Δυσεπίλυτα Αλγοριθμικά Προβλήματα

Ενότητα 1: Nearest neighbors and Clustering

Γιάννης Εμίρης

Τμήμα Πληροφορικής & Τηλεπικοινωνιών
Πανεπιστήμιο Αθηνών

October 10, 2021

Outline

- 1 Similarity search
- 2 Locality sensitive hashing
 - Euclidean space
- 3 Randomized projections
- 4 Centroid-based Clustering
 - Vector spaces
- 5 General Improvements
 - Initialization
 - Reverse assignment
- 6 Evaluation

Exact and Approximate Nearest Neighbors


A d -dimensional space D is equipped with distance function $\text{dist}(\cdot, \cdot)$. Given pointset $P \subset D$, for any query point $q \in D$, its **exact NN** is any point $p_0 \in P$ s.t.:

$$\text{dist}(p_0, q) \leq \text{dist}(p, q), \quad \forall p \in P.$$

For fixed $P \subset D$, and approximation factor $1 > \epsilon > 0$, given a query point q , an ϵ -NN, or **Approximate NN**, of q is any point $p_0 \in P$:

$$\text{dist}(p_0, q) \leq (1 + \epsilon) \cdot \text{dist}(p, q), \quad \forall p \in P.$$

Exact NN is practically linear in n ,
while ANN is sublinear or logarithmic in n



Near neighbors

Approximation factor $c = 1 + \epsilon > 1$.

Definition

Given: finite set of points $P \subset D$, approximation factor $c > 1$, radius r .

Input: query point $q \in D$.

(r, c) -Near Neighbor Decision problem with witness:

If $\exists p_0$ within radius r , output any $p : \text{dist}(q, p) \leq c \cdot r$;

If $\nexists p$ within radius cr , then report NULL;

otherwise output any point p within cr or NULL.

(r, c) -Near Neighbor Range search: Report any subset of

*$\{p \in P : \text{dist}(q, p) \leq c \cdot r\}$ that includes **all** $\{p \in P : \text{dist}(q, p) \leq r\}$.*

In practice, if c is not given, take $c = 1$.

Distances in vector spaces


Definition

The family of ℓ_k norms, for vectors $x, y \in \mathbb{R}^d$, defines various distances:

$$\text{dist}_{\ell_k}(x, y) = \|x - y\|_k = \sqrt[k]{\sum_{i=1}^d |x_i - y_i|^k}, \quad k \geq 0.$$

Examples:

- ℓ_2 : Euclidean distance,
- ℓ_1 : Manhattan (taxicab) distance,
- ℓ_∞ : max distance.
- ℓ_0 : $\|v\|_0 = \# \text{nonzero entries in } v$.



Distance Measure

Definition (Metric)

A distance measure $d : D^2 \rightarrow \mathbb{R}$ is a function that satisfies:

- Non-negativity: $d(x, y) \geq 0$
- Isolation: $x \neq y \Leftrightarrow d(x, y) > 0$
- Symmetry: $d(x, y) = d(y, x)$
- Triangle inequality: $d(x, y) \leq d(x, z) + d(z, y)$

It follows that $d(x, x) = 0$, and $|d(x, z) - d(z, y)| \leq d(x, y)$.

Distances in vector spaces (e.g. ℓ_k , Hamming) are distance measures. Thanks to their (vector) representation, we can also compute attributes, s.t. the mean of a set, or its total order.

NN in \mathbb{R}

Sort/store the points (in balanced binary search tree), use binary search for queries, then:

- Preprocessing in $O(n \log n)$ time
- Data structure requiring $O(n)$ space
- Answer the query in $O(\log n)$ time


A hash-table with M buckets offers a solution with

- preprocessing in $O(M + n) = O(n)$ time
- space $O(M + n) = O(n)$
- query time $O(1)$

assuming constant time for hashing and constant number of items per bucket.

NN in \mathbb{R}^2

- Preprocessing: Voronoi Diagram in $O(n \log n)$.
- Storage = $O(n)$.
- Given query q , find the cell it belongs (point location) in $O(\log n)$.
NN = site of cell containing q .



NN in \mathbb{R}^d

Exact NN:

- Voronoi diagram uses space/preprocessing = $O(n^{\lceil d/2 \rceil})$; point location methods hardly extend to high dimensions.
- State of the art: kd-trees: $S_p = O(dn)$, Query $\simeq O(d \cdot n^{1-1/d})$, tends to linear in n for large d .
- Randomized (Clarkson'88): $S = O(n^{\lceil d/2 \rceil + \delta})$, $Q = O(2^d \log n)$.

Hence the **Curse of Dimensionality**:

Can we solve NN in poly-time in d and faster than linear-time in n ?

Approximate nearest neighbor (ANN)

- Tree-based data structures: kd-, BBD-trees



(Arya, Mount et al. '98)

- Space = $O(dn)$
- Query = $O((1/\epsilon)^d \log n)$

... still plagued by the curse of dimensionality

- LSH (Indyk, Motwani '98) (Andoni, Indyk et al. '08)


- Space = $O(dn^{1+\rho})$,
- Query = $O(dn^\rho)$, $\rho = 1/(1+\epsilon)^2 < 1$.

... "beats the curse of dimensionality" (Chazelle)

- Dimensionality reduction (E, Psaros et al. '15-19)

- Space = $O^*(dn)$,
- Query = $O^*(dn^\rho)$, $\rho = 1 - \Theta(\epsilon^2)$

... beats the curse in optimal space



Outline

- 1 Similarity search
- 2 Locality sensitive hashing**
 - Euclidean space
- 3 Randomized projections
- 4 Centroid-based Clustering
 - Vector spaces
- 5 General Improvements
 - Initialization
 - Reverse assignment
- 6 Evaluation

LSH definition


Idea: use hash-table for proximity query, mapping nearby vectors to same bucket

LSH Family

Let $r_1 < r_2$, probabilities $P_1 > P_2$. A family H of functions is (r_1, r_2, P_1, P_2) -sensitive if, for any points $p \neq q$ and any randomly selected function $h \in_R H$,

- if $\text{dist}(p, q) \leq r_1$, then $\text{prob}[h(q) = h(p)] \geq P_1$,
- if $\text{dist}(p, q) \geq r_2$, then $\text{prob}[h(q) = h(p)] \leq P_2$.

$h \in_R H$ is chosen uniformly at random from H .



Construction

Preprocess

- Having defined hash-function family H we can construct several $h_j \in_R H$
- Randomly specify L amplified hash-functions g_1, \dots, g_L , each using k h_j 's
- Store all points to i -th (1-dim) hash-table using $g_i, i = 1, \dots, L$


Figure shows intuition behind h_j 's:

Practical choices:

$L = 5$ (or 6),

HashTable size $n/8$ (or $n/16$),

aiming at $\Theta(1)$ points per bucket.



NN search

Approximate kNN

Input: query q

Let $b \leftarrow \text{Null}$; $d_b \leftarrow \infty$; initialize k best candidates and distances;

for i from 1 to L **do**

for each item p in bucket $g_i(q)$ **do**

if $\text{dist}(q, p) < d_b = k\text{-th best distance}$ **then** $b \leftarrow p$; $d_b \leftarrow \text{dist}(q, p)$

end if

if large number of retrieved items (e.g. $> 10L$) **then return** b // optional

end if

end for

return b ; k best candidates;

end for

For $k > 1$ must maintain k best candidate distances sorted so as to always replace k -th farthest candidate neighbor.

Approximate Range Search

Approximate (r, c) Range search

Input: r , query q

for i from 1 to L **do**

for each item p in bucket $g_i(q)$ **do**

if $\text{dist}(q, p) < r$ **then output** p

end if

if large number of retrieved items (e.g. $> 20L$) **then return** // optional

end if

end for

end for

return

This may miss points within radius r .

Amplification

Hash-tables

LSH creates hash-tables using **amplified** index function g by combining k functions $h_i \in_R H$, chosen uniformly at random with repetition from H . So some h_i may be chosen more than once for some g , or for different g 's.

Notice both h_i, g map objects to integers (indices in the hash-table).

Setting $g(p) = [h_1(p), \dots, h_k(p)]$ defines k -dimensional table with many empty buckets. So, we apply concatenation

$$g(p) = [h_1(p)|h_2(p)|\dots|h_k(p)],$$

or random linear combination, for some large $M \in \mathbb{N}$, random $r_i \in \mathbb{N}$:

$$g(p) = \sum_{i=1}^k r_i h_i(p) \bmod M \in [0, \dots, M).$$

Practical choice $k = 4$ to 6 .

Known LSH-able metrics

- Hamming distance,
- ℓ_2 (Euclidean) distance,
- ℓ_1 (Manhattan) distance,
- ℓ_k distance for any $k \in (1, 2)$,
- ℓ_2 distance on a sphere,
- Cosine similarity,
- Jaccard coefficient.

Recall ℓ_k distance:
$$\text{dist}_{\ell_k}(x, y) = \sqrt[k]{\sum_{i=1}^d |x_i - y_i|^k}.$$

Outline

- 1 Similarity search
- 2 Locality sensitive hashing
 - Euclidean space
- 3 Randomized projections
- 4 Centroid-based Clustering
 - Vector spaces
- 5 General Improvements
 - Initialization
 - Reverse assignment
- 6 Evaluation

Euclidean Space

Recall: $\text{dist}_{\ell_2}(x, y)^2 = \sum_{i=1}^d (x_i - y_i)^2$.

Definition

Let d -vector $v \sim \mathcal{N}(0, 1)^d$ have coordinates identically independently distributed (i.i.d.) by the standard normal (next slide).

Set "window" $w \in \mathbb{N}^*$ for the entire algorithm, pick single-precision real t uniformly $\in_R [0, w)$. For point $p \in \mathbb{R}^d$, define:

$$h(p) = \left\lfloor \frac{p \cdot v + t}{w} \right\rfloor \in \mathbb{Z}.$$


- Essentially project p on the line of v , shift by t , partition into cells of length w .
- In general, $w = 4$ is good but should increase for range queries with large r .
- Also $k = 4$ (but can go up to 10), and L may be 5 (up to 30).

Normal distribution

Vector $v \sim \mathcal{N}(0, 1)^d$ has single-precision real coordinates distributed according to the standard normal (Gaussian) distribution:

$$v_i \sim \mathcal{N}(0, 1), \quad i = 1, 2, \dots, d,$$

with mean $\mu = 0$, variance $\sigma^2 = 1$ (σ is the standard deviation).



The bell curve:

Hash-table

To avoid many empty buckets we **define** ϕ as random combination of the h_i 's:

Classic hash-function

Build 1-dim hash-table with classic index:

$$\phi(p) = [(r_1 h_1(p) + r_2 h_2(p) + \cdots + r_k h_k(p)) \bmod M] \bmod \text{TableSize},$$

s.t. $\text{int } r_i \in_R \mathbb{Z}$, prime $M = 2^{32} - 5$ if $h_i(p)$ are int , $\text{TableSize} = n/4$ (or $n/8$).

ϕ computed in int arithmetic, if all $h_i(p)$, r_i are int (≤ 32 bits).

Can have smaller $\text{TableSize} = n/8$ or $n/16$ (heuristic choice).

Recall $(a \square b) \bmod m = ((a \bmod m) \square (b \bmod m)) \bmod m$

Querying trick

Store object ID along with pointer to object, for all bucket elements.

Object ID

For every p , store

$$\text{ID}(p) = r_1 h_1(p) + r_2 h_2(p) + \dots + r_k h_k(p) \bmod M.$$

Then indexing hash-function is $\phi(p) = \text{ID}(p) \bmod \text{TableSize}$.

ID is locality sensitive: depends on w -length cells on the v -lines.

To avoid computing Euclidean distance to all elements in the bucket, do it only for p : $\text{ID}(p) = \text{ID}(q)$, assuming such p exists.

Complexity

Theorem

For $P \subset \mathbb{R}^d$, $n = |P|$, the hashtable for $(r, 1 + \epsilon)$ -Neighbors offers, whp,
queries in $O(dn^{\frac{1}{1+\epsilon}} \log n)$, space in $O(dn + n^{1+\frac{1}{1+\epsilon}}) \log n$.

Theorem

We solve $(1 + \epsilon)$ -ANN by performing $O(\log \frac{n}{\epsilon})$ many queries to $(r, 1 + \epsilon)$ -Neighbor structures. Binary search is not sufficient.

Corollary

$(1 + \epsilon)$ -ANN performs queries in $O(dn^{\frac{1}{1+\epsilon}} \log n \log \frac{n}{\epsilon})$ using space in $O(dn + n^{1+\frac{1}{1+\epsilon}} \log^3(\frac{n}{\epsilon})/\epsilon^2)$.

Outline

- 1 Similarity search
- 2 Locality sensitive hashing
 - Euclidean space
- 3 Randomized projections**
- 4 Centroid-based Clustering
 - Vector spaces
- 5 General Improvements
 - Initialization
 - Reverse assignment
- 6 Evaluation

Randomized projection

Lemma (Johnson,Lindenstrauss'82)

Given $P \subset \mathbb{R}^d$, $|P| = n$, $\epsilon > 0$, \exists randomized linear $f : \mathbb{R}^d \rightarrow \mathbb{R}^{d'}$, for

$$d' = O(\log n / \epsilon^2),$$

s.t., whp, $(1 - \epsilon)\|p - q\|_2 \leq \|f(p) - f(q)\|_2 \leq (1 + \epsilon)\|p - q\|_2$.


For randomized $d' \times d$ matrix M , projected vectors are Mp for input $p \in \mathbb{R}^d$.

However, the JL Lemma cannot remedy the curse of dimensionality for ANN.

Binary (0/1) hypercube

Projection

- Input: Metric space admitting family of LSH functions h_i .
- Let f_i s.t: for each h_i , $f_i(h_i)$ maps buckets to $\{0, 1\}$ uniformly.
- Preprocess: Store points $p \mapsto [f_1(h_1(p)), \dots, f_{d'}(h_{d'}(p))] \in \{0, 1\}^{d'}$, so data at hypercube vertices, dimension $d' = \lfloor \lg n \rfloor - 1$ to $\lfloor \lg n \rfloor - 3$.



Search

- 1 Project query point to corresponding hypercube vertex.
- 2 Check points in same vertex and nearby vertices in increasing Hamming distance (=1, then 2, etc), until some threshold reached (next slide).
- 3 ANN returns closest candidate, range search all p within query ball.

Complexity

Theorem

For ℓ_1 and ℓ_2 metrics, a (single) hypercube yields an efficient solution for the ϵ -ANN problem with space and preprocessing in $O^*(dn)$, query time in $O^*(dn^\rho)$, $\rho = 1 - \Theta(\epsilon^2)$. The data structure succeeds with constant probability.

Implementation Parameters

- d' : larger implies finer mapping so search can stop earlier (lower thresholds); increases storage and preprocessing.
- Thresholds: #points to be checked in \mathbb{R}^d , max #vertices probed, bound on Hamming distance of vertices probed.

The Hamming distance of strings $x, y \in \{0, 1\}^k$ equals # different bits in x, y . Here the vertices of the (unique) hypercube are the hash-table buckets.

Outline

- 1 Similarity search
- 2 Locality sensitive hashing
 - Euclidean space
- 3 Randomized projections
- 4 Centroid-based Clustering**
 - Vector spaces
- 5 General Improvements
 - Initialization
 - Reverse assignment
- 6 Evaluation


Clustering

Definition (k clusters)

Given n objects (and $k > 1$) partition the objects into k clusters so as to optimize some objective function.

- Objects in the same cluster are more "similar" (or closer) to each other than to those in other clusters.
- Possible criteria: minimizing the total distance among all cluster points, minimizing the distance of cluster points to some center, etc.
- Variations: k is unknown and computed, e.g., by the Silhouette method.

Good Clustering, with centroids



Approaches

- hierarchical (agglomerative): each point initializes a cluster, merge until stopping criterion, e.g., predetermined number of clusters, or if merging creates cluster with points too far apart.
- centroid-based (point-assignment): given some initial clusters (with centroids), assign points to "best" cluster (nearest centroid); redefine centroids and repeat. Might allow combining / splitting clusters. Example: k-means (our focus).

(Ullman et al: Mining Massive datasets)

(Theodoridis-Koutroumbas: Pattern Recognition)

Outline

- 1 Similarity search
- 2 Locality sensitive hashing
 - Euclidean space
- 3 Randomized projections
- 4 Centroid-based Clustering**
 - Vector spaces**
- 5 General Improvements
 - Initialization
 - Reverse assignment
- 6 Evaluation


Vector spaces

Problem definition

- Clustering that minimizes objective function.
- k is given.
- Centroids may **not** be in the input; computed per cluster.

k-means

- **k-means** is the most common problem: Main algorithms:
 - Lloyd's algorithm is standard.
 - Elkan's uses triangular inequality to accelerate updates.
- In Euclidean space, assignment is point location to k Voronoi cells.



Lloyd's for k-means

Algorithm (EM)

Initialize k centers randomly (or using some strategy).

- 1 Assignment (Expectation): Assign each object to its nearest center.
- 2 Update (Maximization): Calculate mean $\frac{1}{T} \sum_{i=1}^T \vec{v}_i$ per cluster, make it new center, where $T = \# \text{objects } \vec{v}_i \text{ in cluster}$.

Repeat the two steps until there is no, or little, change in the assignments.

Properties

- Each distance calculation = $O(d)$ because vectors in \mathbb{R}^d .
- Assignment = $O(nkd)$, Update = $O(nd)$,
- #iterations unknown, in practice $\ll n$.
- Converges to local minimum in Euclidean space (depends on initialization)

k-means: Objective function

Typical ambient space is \mathbb{R}^d but can generalize to metric space \mathcal{Z} .

Minimization function

In any metric space over vectors \mathcal{Z} with distance metric d , let the dataset be $X = \{x_1, \dots, x_n\} \subseteq \mathcal{X} \subseteq \mathcal{Z}$, and $k > 1$. Given centroids $C \subset \mathcal{Z}$, let

$$d(x_i, C) = \min_{c \in C} d(x_i, c).$$

Consider vector $v(C) = [d(x_1, C), \dots, d(x_n, C)]$. The k-means objective is:

$$\min_{C \subseteq \mathcal{Z}, |C|=k} \|v(C)\|_2^2 = \min_{C \subseteq \mathcal{Z}, |C|=k} \sum_{i=1}^n d(x_i, C)^2.$$

The k -means objective is NP-hard to check, but for the ℓ_2 metric, Lloyd's algorithm converges "quickly" to a *local* minimum.

Variations

Recall $X = \{x_i\}$ is input, $v(C) = [d(x_1, C), \dots, d(x_n, C)]$, where $C \subset \mathcal{Z}$ are the k centroids in the ambient space \mathcal{Z} . The **k-median** objective is:

$$\min_{C \subseteq \mathcal{Z}, |C|=k} \|v(C)\|_1 = \min_{C \subseteq \mathcal{Z}, |C|=k} \sum_{i=1}^n d_{\ell_1}(x_i, C).$$

Further objective functions exist when the centroids belong to the dataset:

k-medoid: $\min_{C \subseteq X, |C|=k} \|v(C)\|_1$.

k-center: $\min_{C \subseteq X, |C|=k} \|v(C)\|_\infty$.

k-medians clustering

Lloyd's Algorithm (EM)


Initialize k centers.

- 1 Expectation: Assignment of vectors to nearest center in ℓ_1 metric.
- 2 Maximization: Update center by (marginal) median $c \in \mathbb{R}^d$ defined s.t. c_j is **median** of $\{v_{1j}, v_{2j}, \dots, v_{Tj}\}$, $1 \leq j \leq d$, T = cluster cardinality.


Repeat the two steps until there is no change in the assignments.

Recall the median of n sorted values is the $\lceil n/2 \rceil$ -th largest value.


Random initialization $k = 4$




Lloyd's 1




Lloyd's 2




Lloyd's 3



Lloyd's 4



Fails to find 4 clusters



Outline

- 1 Similarity search
- 2 Locality sensitive hashing
 - Euclidean space
- 3 Randomized projections
- 4 Centroid-based Clustering
 - Vector spaces
- 5 General Improvements**
 - Initialization
 - Reverse assignment
- 6 Evaluation

Outline

- 1 Similarity search
- 2 Locality sensitive hashing
 - Euclidean space
- 3 Randomized projections
- 4 Centroid-based Clustering
 - Vector spaces
- 5 General Improvements**
 - **Initialization**
 - Reverse assignment
- 6 Evaluation

Improve Initialization 1: Spread-out

INITIALIZATION++ : K-MEANS++ / K-MEDOIDS++:

- (1) Choose a centroid uniformly at random; $t \leftarrow 1$.
- (2) \forall non-centroid point $i = 1, \dots, n - t$,
let $D(i) \leftarrow$ min distance to some centroid, among t chosen centroids.
- (3) Choose new centroid: r chosen with probability proportional to $D(r)^2$:

$$\text{prob}[\text{choose } r] = D(r)^2 / \sum_{i=1}^{n-t} D(i)^2.$$

Let $t \leftarrow t + 1$.

- (4) Go to (2) until $t = k = \text{given \#centroids}$.

Expected approximation ratio = $O(\log k)$ (Arthur-Vassilvitskii:SODA'07)

Similar algo for 2-approx of k -center (NP-hard prob)

Implement initialization++

Given $D(i) > 0$, $i = 1, \dots, n - t$, compute $n - t$ (`float`) partial sums

$$P(r) = \sum_{i=1}^r D(i)^2, \quad r = 1, \dots, n - t,$$

and store them in binary tree (or array) P . To avoid the $P(r)$'s being very large, we may normalize all $D(i)$'s by dividing them by $\max_i D(i)$.

Pick a uniformly distributed `float` $x \in [0, P(n - t)]$ and return

$$r \in \{1, 2, \dots, n - t\} : P(r - 1) < x \leq P(r),$$

where $P(0) = 0$; r chosen with probability proportional to $P(r) - P(r - 1) \sim D(i)^2$. Can find r by binary search in P .

Improve Initialization 2: Concentrate

Select centroids close to dataset's center of mass (and to each other) as follows.

(1) Calculate symmetric $n \times n$ distance matrix of all objects, i.e. all distances d_{ij} from every object $i = 1, \dots, n$ to every other object $j = 1, \dots, n, i \neq j$.

(2) For object i compute

$$v_i = \sum_{j=1}^n \frac{d_{ij}}{\sum_{t=1}^n d_{jt}}, \quad i = 1, \dots, n.$$

(3) Return the k objects with k smallest v_i values.

Algorithm proposed in (Park-Jun'09).

Outline

- 1 Similarity search
- 2 Locality sensitive hashing
 - Euclidean space
- 3 Randomized projections
- 4 Centroid-based Clustering
 - Vector spaces
- 5 General Improvements**
 - Initialization
 - Reverse assignment**
- 6 Evaluation

Assignment by direct method

Exact approach -- for little data

At each iteration (Lloyd's):

- 1 For every point, compute distance to every centroid.
- 2 Return (exact) nearest centroid.

Approximate approach -- for big data

At each iteration:

- 1 Index k centroids into data-structure, e.g. LSH hashtables.
- 2 For every non-centroid point, run ANN to find nearest centroid.
- 3 Return (approximate) nearest centroid.

This is the standard approach in almost all big data implementations today.


Assignment by Range search

Reverse approach (ANN)

- Index n points into L hashtables: once for entire algorithm.
- LSH TableSize $\leq n/8$: avoid buckets with very few items.
- At each iteration, for each centroid c , range/ball queries centered at c .
- Mark assigned points: either move them at end of LSH buckets (and insert "barrier", or mark them using "flag" field).
- Multiply radii by 2, start with $\min(\text{dist between centers})/2$; centers mapped to buckets once. Until most balls get no new point.
- For a given radius, if a point lies in ≥ 2 balls, compare its distances to the respective centroids, assign to closest centroid.
- At end: for every unassigned point, compare its distances to all centroids

Reverse assignment: example

Double radii until all points assigned, or most balls get no new point, or radius reaches some threshold.



Outline

- 1 Similarity search
- 2 Locality sensitive hashing
 - Euclidean space
- 3 Randomized projections
- 4 Centroid-based Clustering
 - Vector spaces
- 5 General Improvements
 - Initialization
 - Reverse assignment
- 6 Evaluation

Silhouette

- For $1 \leq i \leq n$, $a(i)$ = average ℓ_1 distance of i to objects in same cluster
- Let $b(i)$ = average ℓ_1 distance of i to objects in *next best* (neighbor) cluster, i.e. cluster of 2nd closest centroid.

Silhouette of Object i

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}} = \begin{cases} 1 - a(i)/b(i), & \text{if } a(i) < b(i) \\ 0, & \text{if } a(i) = b(i) \\ b(i)/a(i) - 1, & \text{if } a(i) > b(i) \end{cases} \in [-1, 1].$$

Interpret silhouette

- $s(i) \rightarrow 1$: i seems correctly assigned to its cluster;
- $s(i) \simeq 0$: borderline assignment (but not worth to change);
- $s(i) \rightarrow -1$: i would be better if assigned to next best cluster.

Silhouette: Cluster and clustering

Specific clusters

- Evaluate a cluster: Compute average $s(i)$ over all i in some cluster.
- If k is too large or too small, some clusters shall display much smaller silhouettes than the rest.
- Silhouette plots are used to improve k : try different k 's and see if clusters have roughly equal silhouettes.

Overall Clustering

Overall Silhouette coefficient = average $s(i)$ over $i = 1, \dots, n$.

High if well clustered, low may indicate bad k (or existence of outlier points).